

# train\_cnn

October 26, 2019

## 1 Training a Texture Classification CNN on MVTec AD

### 1.1 Imports

```
In [1]: import os
import numpy as np
import tensorflow
import keras

from keras.preprocessing.image import ImageDataGenerator
from keras.models import Sequential
from keras.layers import Conv2D, MaxPooling2D, GlobalAveragePooling2D, Dense
from keras import regularizers
from keras.models import load_model

# from tensorflow.python.summary import summary_iterator
from tensorflow.core.util import event_pb2
from tensorflow.python.lib.io import tf_record

import matplotlib.pyplot as plt
```

Using TensorFlow backend.

Find the Python distribution used for the Jupyter notebook:

```
In [2]: import sys
print(sys.executable)
```

C:\Users\chanlynd\Anaconda3\python.exe

If Keras/TensorFlow are not yet installed in below Python distribution, run `pip install keras/pip install tensorflow`.

### 1.2 Load data

```
In [3]: # Settings
DATASET_ROOT = 'data/train'
```

```

CLASSES = ['carpet', 'grid', 'leather', 'tile', 'wood']

# Read in the data
train_datagen = ImageDataGenerator(rescale=1./255,
                                   shear_range=0.2,
                                   rotation_range=30,
                                   horizontal_flip=True,
                                   vertical_flip=True)

train_generator = train_datagen.flow_from_directory(DATASET_ROOT,
                                                    class_mode='categorical',
                                                    interpolation='bilinear',
                                                    target_size=(224, 224),
                                                    batch_size=16,
                                                    shuffle=True,
                                                    classes=CLASSES)

unique, counts = np.unique(train_generator.labels, return_counts=True)
print(dict(zip(CLASSES, counts)))

```

Found 1266 images belonging to 5 classes.

```
{'carpet': 280, 'grid': 264, 'leather': 245, 'tile': 230, 'wood': 247}
```

### 1.3 Define model architecture

In [4]: *# Read in the model*

```

weight_decay = 5e-4
model = Sequential()
model.add(Conv2D(64, kernel_size=(3, 3), padding='same', activation='relu', input_shape=(224, 224, 3)))
model.add(Conv2D(64, kernel_size=(3, 3), padding='same', activation='relu', kernel_regularizer=regularizers.l2(weight_decay)))
model.add(MaxPooling2D(pool_size=(2, 2)))

model.add(Conv2D(128, kernel_size=(3, 3), padding='same', activation='relu', kernel_regularizer=regularizers.l2(weight_decay)))
model.add(Conv2D(128, kernel_size=(3, 3), padding='same', activation='relu', kernel_regularizer=regularizers.l2(weight_decay)))
model.add(MaxPooling2D(pool_size=(2, 2)))

model.add(Conv2D(256, kernel_size=(3, 3), padding='same', activation='relu', kernel_regularizer=regularizers.l2(weight_decay)))
model.add(Conv2D(256, kernel_size=(3, 3), padding='same', activation='relu', kernel_regularizer=regularizers.l2(weight_decay)))
model.add(Conv2D(256, kernel_size=(3, 3), padding='same', activation='relu', kernel_regularizer=regularizers.l2(weight_decay)))

model.add(GlobalAveragePooling2D())
model.add(Dense(len(CLASSES), activation='softmax', kernel_initializer=keras.initializers.RandomNormal(mean=0.0, stddev=0.05),
                                                         bias_initializer=keras.initializers.Zeros(), kernel_regularizer=regularizers.l2(weight_decay)))

# Compile the model
model.compile(loss=keras.losses.categorical_crossentropy,
              optimizer=keras.optimizers.Adadelta(),
              metrics=['accuracy'])

```

## 1.4 Train the model

WARNING: this will train the model for 100 epochs, do not run unless necessary

```
In [ ]: # def lr_scheduler(epoch):
#         # return 1e-3 * (.5 ** (epoch // 5))
#         return 1e-2 * (.5 ** (epoch // 20))
#     lr_reduce_cb = keras.callbacks.LearningRateScheduler(lr_scheduler)
#     tensorboard_cb = keras.callbacks.TensorBoard(log_dir='log', write_graph=True)
#     model.fit_generator(generator=train_generator,
#                         steps_per_epoch=train_generator.n // 16,
#                         epochs=100,
#                         callbacks=[lr_reduce_cb, tensorboard_cb],
#                         verbose=2)
```

## 1.5 Save the model

WARNING: this will overwrite the saved model provided, do not run unless necessary

```
In [ ]: # if not os.path.exists(save_dir):
#         os.makedirs(save_dir)
#     model.save('save/texturenet.h5')
```

## 1.6 Load the saved model

```
In [5]: model = load_model('save/texturenet.h5')
```

## 1.7 Load training progress from Tensorboard

```
In [6]: def my_summary_iterator(path):
#         for r in tf_record.tf_record_iterator(path):
#             yield event_pb2.Event.FromString(r)

train_acc = []
train_loss = []
train_lr = []
print('log/' + os.listdir('log')[0])
for e in my_summary_iterator('log/' + os.listdir('log')[0]):
    for v in e.summary.value:
        if v.tag == 'acc':
            train_acc.append(v.simple_value)
        elif v.tag == 'loss':
            train_loss.append(v.simple_value)
        elif v.tag == 'lr':
            train_lr.append(v.simple_value)
```

log/events.out.tfevents.1571952078.HADES

WARNING:tensorflow:From <ipython-input-6-00d5ea4b2a47>:2: tf\_record\_iterator (from tensorflow.


Instructions for updating:

Use eager execution and:  
`tf.data.TFRecordDataset(path)`

## 1.8 Plot training accuracy

```
In [7]: plt.plot(np.arange(len(train_acc)), train_acc)
        plt.xlabel('Epoch')
        plt.ylabel('Training Accuracy')
```

```
Out[7]: Text(0, 0.5, 'Training Accuracy')
```




```
In [8]: print('Final training accuracy=%f' % train_acc[-1])
```

```
Final training accuracy=0.993600
```

```
In [9]: plt.plot(np.arange(len(train_loss)), train_loss)
        plt.xlabel('Epoch')
        plt.ylabel('Training Loss')
```

```
Out[9]: Text(0, 0.5, 'Training Loss')
```




```
In [10]: print('Final training loss=%f' % train_loss[-1])
```

```
Final training loss=0.635351
```

```
In [11]: plt.plot(np.arange(len(train_lr)), train_lr)
plt.xlabel('Epoch')
plt.ylabel('Training Learning Rate')
```

```
Out[11]: Text(0, 0.5, 'Training Learning Rate')
```



In [ ]: